

CAFE: A Design-of-Experiments Platform for Evaluating Compound AI Systems

Fabian Lukassen¹ Christoph Weisser² Alexander Silbersdorff¹

¹University of Göttingen ²Bielefeld University of Applied Sciences and Arts (HSBI)

fabian.lukassen@stud.uni-goettingen.de christoph.weisser@hsbi.de
asilbersdorff@uni-goettingen.de

Abstract

Once an AI application leaves the single-model setting, it becomes a *compound AI system* (Zaharia et al., 2024): a pipeline of interacting, often non-differentiable techniques—retrieval, reranking, context assembly, prompting, one or more model calls, tool use, and guardrails—any of which can be swapped. Practitioners who build such systems face a combinatorial configuration space and three coupled questions that current tooling answers only indirectly: *which techniques drive answer quality, what is the best configuration, and where is effort best spent*—for instance, on prompt engineering versus a stronger retriever? We present **CAFE** (Compound-AI Factorial Evaluation), an open-source platform that brings *design of experiments* (DoE) to compound AI evaluation. CAFE treats any technique registered in the pipeline as an experimental *factor*, generates full or fractional factorial designs over the chosen factors, executes the resulting configurations against the system as a black box with replicated runs, and collects ordinal quality judgments from a configurable LLM judge and from human domain experts on a shared rubric. A statistics layer attributes answer-quality variance to techniques and their interactions using cumulative link mixed models, identifies the best configuration and the highest-leverage factors, reports effect sizes with uncertainty, and quantifies judge–human agreement. We describe the system, a live deployment evaluating a retrieval-augmented legal-research assistant for German works councils with raters from partner organisations, and a programme for the platform’s empirical validation. CAFE is a Docker-deployable web application with a public demo.

1 Introduction

As soon as an AI application moves beyond a single model call, it becomes what the literature now calls a *compound AI system*: an application that solves a

task with several interacting techniques rather than one monolithic network (Zaharia et al., 2024; Chen et al., 2025a). The techniques are open-ended—retrieval, reranking, context assembly, prompting, one or more model calls, tool and web-search use, and input/output guardrails are common examples, but the set is not fixed and grows as practitioners invent new ones. Two properties make these systems hard to reason about. They are *combinatorial*: each technique can be realised in several ways, and the configurations multiply. And they are typically *non-differentiable*—they contain search engines, retrievers, and tool calls with no gradient—so they cannot be tuned by backpropagation and must be treated as black boxes (Zaharia et al., 2024; Lee et al., 2025). A substantial literature has grown around *building* and *optimizing* such pipelines, from declarative programming models (Khattab et al., 2024; Yuksekgonul et al., 2025) to model selection (Chen et al., 2025b), system-level alignment (Wu et al., 2026), and end-to-end configuration search (Kartal et al., 2025; Barker and Bell, 2025).

Evaluating these systems, however, has lagged behind building them. A practitioner who has assembled a retrieval-augmented assistant and can choose among, say, three base models, two retrieval corpora, and three prompt templates is confronted with eighteen candidate configurations and a cluster of coupled questions that available tools answer only indirectly: *which techniques drive answer quality, and by how much? Which configuration is best, and is the difference I observe real? And where should I spend effort next—is it worth engineering the prompt, or would a stronger retriever pay off more?* The last question matters in practice: engineering time is finite, and a quantitative apportioning of quality to techniques tells a team where its next hour is best invested. These questions are harder than they look, for three reasons that recur throughout the recent literature.

First, components **interact**, often non-

monotonically. Upgrading a retriever can *lower* end-to-end quality if the downstream generator is sensitive to how evidence is formatted, and adding model calls can help easy queries while hurting hard ones (Chen et al., 2024). Evaluating one factor at a time therefore mismeasures the system, because it cannot estimate interactions.

Second, large language models are **nondeterministic** even under nominally deterministic settings. Repeated decoding at temperature zero yields accuracy swings of up to fifteen points across runs (Atil et al., 2024), agentic pass@1 estimates vary by several points on identical inputs (Bjarnason et al., 2026), and LLM judges disagree with themselves across repeats (Haldar et al., 2025). A comparison of two configurations from a single run each is, in consequence, statistically invalid.

Third, the quality signals used to compare configurations are typically **ordinal**—a rubric level such as *worse/on par/better*, or a Likert rating—yet are routinely analysed as if they were interval-scaled means. Treating ordinal ratings as interval data is known to underpower NLP evaluations, sometimes severely (Howcroft and Rieser, 2021; Taylor et al., 2023).

These three problems—interactions, nondeterminism, and ordinal measurement—are exactly the problems that *design of experiments* (DoE) and mixed-effects modelling were developed to solve in the empirical sciences. A factorial design estimates main effects *and* interactions from a balanced, efficient set of runs; replication and random effects separate signal from run-to-run noise; and ordinal regression draws valid conclusions from rubric data. Our contribution is to package this methodology into a usable system for the people who build compound AI pipelines.

We present **CAFE** (Compound-AI Factorial Evaluation), an open-source platform with four contributions:

1. A **pipeline-as-factors** abstraction. Any technique in a compound system—a retriever, a corpus, a reranker, a prompt, a model, a tool, a guardrail—is registered in code with a one-line decorator and surfaced automatically in the interface as a *factor* with selectable *levels*; the set of testable techniques is open and extends with the codebase, so that defining an experiment requires no statistical programming.
2. A **DoE engine** that generates full and frac-

tional factorial designs over the chosen factors, executes the configurations against the target system as a black box, and replicates runs to capture nondeterminism.

3. A **configurable dual-rating** workflow in which the same answers are scored on a shared ordinal rubric by one or more LLM judges—each a freely chosen model, prompt, rubric, and repetition count—and by human domain experts, with inter-rater reliability reported between them.
4. A **variance-attribution** analysis layer that fits cumulative link mixed models with question-level random effects, reports per-factor effect sizes and uncertainty, identifies the best configuration and the highest-leverage factors, and serves the results through a researcher-facing dashboard.

CAFE is built for the teams who assemble and compare compound pipelines—practitioners and research groups that run evaluations—together with the domain experts who supply the quality judgments.

2 Related Work

Building and optimizing compound systems. The compound-systems framing (Zaharia et al., 2024; Chen et al., 2025a) has produced a family of optimizers: DSPy compiles declarative pipelines and tunes their prompts (Khatab et al., 2024), TextGrad backpropagates textual feedback (Yuksekgonul et al., 2025), LLMSelector allocates models to modules (Chen et al., 2025b), Optimas aligns local rewards with global performance (Wu et al., 2026), and RAGSmith searches tens of thousands of RAG configurations with a genetic algorithm (Kartal et al., 2025). These systems *search* a large space for a winning configuration but return little about *why* it wins: no calibrated per-technique effect sizes, no interaction structure, and no account of whether a margin survives nondeterminism. CAFE is complementary and, for the modest designs practitioners actually compare, answers more: over the configurations in a design it both reports the best configuration *and* explains the space—how much each technique and interaction contributes, with uncertainty—which is what tells a team whether a knob is worth handing to an optimizer at all (Chen et al., 2024). CAFE does not (yet) perform large-scale search; screening with fractional designs is its analogue, and Bayesian

optimization is future work.

Evaluation frameworks. Holistic benchmarks such as HELM (Liang et al., 2023) characterise *models* across scenarios but not *pipelines* across component choices. RAG evaluators—RAGAS (Es et al., 2024) and ARES (Saad-Falcon et al., 2024)—score pipeline components such as faithfulness and context relevance in isolation using LLM judges, and ARES adds prediction-powered confidence intervals (Fisch et al., 2024); neither varies component choices in a designed experiment to attribute end-to-end variance. Behaviour-driven evaluation (Bhagat et al., 2025) generates scenario-based tests that surface more failures than static benchmarks, but does not model which component is responsible. Production platforms (Inspect AI (UK AI Security Institute, 2024), LangSmith (LangChain, 2024), Argilla (Argilla, 2024)) provide execution harnesses, annotation queues, and LLM-as-judge scoring; CAFE occupies the same layer for execution and rating and adds what these tools leave to the user—the experimental design and the statistical analysis that turn a pile of scores into per-technique attribution.

Statistics for AI evaluation. A growing line of work argues that AI evaluation must be treated as statistical inference: quantifying benchmark uncertainty (Wang et al., 2025), modelling nested structure with hierarchical Bayesian models (Luettgau et al., 2025), decomposing measurement error in LLM pipelines (Messing, 2026), and using ordinal rather than interval models for human ratings (Howcroft and Rieser, 2021; Syiem and Velloso, 2026; Taylor et al., 2023). Variance-based sensitivity analysis (Saltelli et al., 2008) is the classical instrument for apportioning output variance to inputs. The works closest to ours are two. Haase et al. (2026) partition output variance into prompt, model, and sampling components, but for standalone creative generation, with *linear* mixed models on an interval quality score from a single automatic rater, and with no retrieval stage; we model the response *ordinally*, attribute variance across the *full pipeline* (including the retrieval corpus), and use both human-expert and LLM-judge ratings. Mustahsan et al. (2025) quantify run-to-run inconsistency with the intraclass correlation coefficient, a single global between-/within-query split; a cumulative link mixed model instead attributes variance to *named pipeline factors and their interactions*, of which run-to-run noise is only one

Listing 1: Registering techniques. The `stage.name` key makes retrieve a factor with `bm25` among its levels; adding a class is the only step needed to make a new technique testable.

```
from cafe.registry import register, Technique

@register("retrieve.bm25") # stage = retrieve
class BM25(Technique): # level = bm25
    def run(self, query, ctx):
        return bm25_search(query, k=ctx.top_k)

@register("prompt.neutral") # stage = prompt
class NeutralPrompt(Technique):
    text = "You are a neutral legal assistant ..."
```

term. More broadly, CAFE’s contribution is the integration—design, ordinal analysis, dual human/judge rating with reported reliability, and explicit replication—into one deployed, reusable platform with a real-world case study, more than any single technique in isolation. Recent factorial studies of LLM outputs (Lukassen et al., 2026; Vazquez et al., 2025) show the design methodology is timely; we provide the infrastructure to apply it routinely.

3 System Overview

CAFE is a web application built around a single premise: the target compound AI system is a *black box* reachable over HTTP, and everything we vary about it can be expressed as experimental factors. Figure 1 sketches the architecture; we describe each layer in turn.

Pipeline as factors. CAFE models the target pipeline as an ordered sequence of *stages*; a stage is any point at which a technique can be swapped—retrieval, reranking, context assembly, prompting, the model call, tool use, or input/output guardrails are typical, but the list is open and defined by the codebase, not fixed by CAFE. A concrete *technique* implements a stage—a particular retriever, a corpus, a system prompt, a model binding, a guardrail—and is registered with a one-line decorator whose key is `stage.name`. All techniques registered for a stage become the selectable *levels* of that stage’s factor and appear automatically in the experiment UI. Listing 1 shows two registrations.

Because registration is the single source of truth, what can be tested is always exactly what is implemented, and extending the platform to a new retrieval method, reranker, or guardrail is a self-contained code change with no edit to the experiment UI; the registered techniques are browsable in the interface (Appendix A, Figure 2).




Figure 1: CAFE architecture. A registry exposes pipeline techniques as factors; the DoE engine generates a design; an execution worker runs each configuration against the target system as a black box; answers are rated by one or more configurable LLM judges and by human experts on a shared rubric; the statistics layer attributes variance and serves dashboards.

Defining and running a study. In the interface an experimenter (i) selects *questions*, each optionally paired with an expert *reference answer*; (ii) picks the factors to vary and ticks the levels of each—e.g. $\text{model} \in \{A, B, C\}$, $\text{retrieve} \in \{\text{bm25}, \text{dense}\}$, $\text{prompt} \in \{\text{neutral}, \text{strict}\}$; and (iii) chooses a design (Figure 3). CAFE reports the resulting cell count and generates the configurations: a full factorial for small designs, or a fractional factorial that reduces the run count while keeping main effects and low-order interactions estimable, with the alias structure shown. This screen-then-refine workflow is standard DoE practice (Saltelli et al., 2008; Vazquez et al., 2025)—screen many factors cheaply, then estimate interactions among the survivors with a full design. An execution worker then claims configurations from a queue and runs each selected question through the target system over HTTP, bounding per-model concurrency to respect rate limits. Each (configuration, question) cell is executed with *replication*, because outputs vary run to run even at temperature zero (Atil et al., 2024; Bjarnason et al., 2026); the repeats let the analysis estimate within-cell variance instead of mistaking it for a treatment effect. Per-stage timings are recorded so latency can be analysed alongside quality.

Rating. Answers are scored on a configurable *rubric* that fixes an ordinal scale—for example a three-level *worse/on par/better* comparison against the reference answer, though the scale, levels, and judge prompt are all user-defined. Two rater types apply the rubric. An *LLM judge* is fully configurable: the experimenter chooses its model, its prompt, the rubric it applies, and how many times it rates each answer, and several judges can be run

side by side; our current deployment uses a single judge prompt, but nothing in the design privileges it. Repeated judge passes are retained as separate ratings, not averaged, so that judge self-inconsistency (Halder et al., 2025) is measured rather than hidden. *Human domain experts* rate through a dedicated, localized interface (Figure 4); access is by invitation so the operating team controls who contributes. Recording both rater types against one scale is what later permits reliability analysis between them.

Analysis and dashboards. When a study completes, the statistics layer (Section 4) fits the variance-attribution model and serves a results dashboard (Figure 5): the relative contribution of each factor, the best level of each factor with significance, the best end-to-end configuration, the verdict distribution, and pipeline-stage timings. Results are normalised across rubrics so that studies using different scales can be compared at the portfolio level.

Deployment and access. The system is packaged with Docker Compose—a FastAPI backend, a PostgreSQL store, the execution worker, a React frontend, and a reverse proxy—and runs on a single modest server. Two roles are supported: an *administrator* role for the operating team, who define studies and manage the catalogue, and *organisation* accounts for external partners, who may submit questions with reference answers and contribute human ratings.

4 Statistical Methodology

CAFE’s analytical core treats a completed study as a designed experiment with an ordinal response. Let y_{ijr} be the rubric verdict for question i under configuration j on replicate r , taking one of K or-

dered levels. Configuration j is described by its factor levels (model, corpus, prompt, and so on). We model the cumulative probabilities of the verdict with a *cumulative link mixed model* (CLMM) (Taylor et al., 2023; Howcroft and Rieser, 2021):

$$\text{logit } \Pr(y_{ijr} \leq k) = \theta_k - (x_j^\top \beta + u_i + b_{ij}), \quad (1)$$

where $\theta_1 < \dots < \theta_{K-1}$ are ordered thresholds, $x_j^\top \beta$ collects the fixed effects of the factors and their interactions, u_i is a random intercept for question i , and b_{ij} is a question $\times$ configuration interaction that lets a configuration’s quality vary by question. Run-to-run nondeterminism enters as the latent residual of the cumulative-link model; crucially, repeated ratings of the same answer (whether replicate executions or repeated LLM-judge passes) enter as separate observations r rather than being averaged, so that the nondeterminism documented by Atil et al. (2024) and Halder et al. (2025) is estimated rather than silently collapsed—averaging ordinal verdicts would reintroduce exactly the interval-treatment problem we avoid. Because the platform collects judgments from two rater types and several human raters, Eq. (1) is fit per rater type and, when human raters are pooled, extended with a rater random intercept. Modelling the response on the ordinal scale rather than averaging level indices avoids the loss of power and the distortions that follow from treating rubric levels as interval data (Howcroft and Rieser, 2021; Syiem and Velloso, 2026); the question random effect u_i accounts for the fact that some questions are simply harder, so that factor effects are estimated within question rather than confounded with question sampling. With few rubric levels we check the proportional-odds assumption and relax it to a partial-proportional-odds model where it fails.

Variance attribution. For each factor we report a standardised effect size together with its contribution to an analysis of deviance (Type-II likelihood-ratio tests), and we summarise the relative importance of factors with a McKelvey–Zavoina-style pseudo- R^2 decomposition on the model’s latent scale, alongside the estimated random-effect variances. We are explicit that this is an approximate decomposition: the latent residual scale of a logit cumulative-link model is fixed rather than estimated, so there is no single canonical “percentage of variance” as in a linear model, and we report the decomposition as a ranking aid rather than an exact partition. This is the variance-based sensitiv-

ity view (Saltelli et al., 2008) adapted, with these caveats, to an ordinal, replicated, designed setting. Where a factor is shown to matter, pairwise contrasts identify its best level. A best end-to-end configuration can be read off the model when the design is full factorial or higher-order interactions are negligible; under a fractional design only main effects and the explicitly modelled low-order interactions are estimable, and CAFE reports the alias structure so that this limitation is visible rather than hidden.

Reliability. Because the quality signal is itself produced by raters, CAFE quantifies its trustworthiness. Inter-rater reliability between the LLM judge and the human experts—and among human experts—is computed with Krippendorff’s α (Krippendorff, 2004), which handles ordinal scales and missing ratings. Reported alongside the variance attribution, α tells a reader whether “best by our ratings” can be trusted as “best.” This complements a growing literature that calibrates and audits LLM judges—surveys of the paradigm (Gu et al., 2024), analyses of judge–human agreement (Han et al., 2025), calibration against sparse human labels with statistical guarantees (Boyeau et al., 2025; Fisch et al., 2024; Park et al., 2025), bias detection (Dubois et al., 2025), and validation under rating indeterminacy (Guerdan et al., 2025)—which CAFE draws on rather than replaces.

Why this matters. The combination is deliberate. Single-run point comparisons over interval-coded means—still common in practice—can manufacture differences that are artefacts of nondeterminism or of scale misuse (Atil et al., 2024; Messing, 2026). Equation (1) instead separates the treatment effects of components from question difficulty and from run-to-run noise, and reports them on the scale the data actually live on. The platform makes this analysis the default rather than an expert afterthought.

5 Deployment

CAFE is deployed to evaluate a retrieval-augmented assistant that supports German works-council members (*Betriebsräte*) with legal research. The assistant is a natural compound system: a question is answered by retrieving from a legal corpus, assembling context, applying a system prompt, and calling one of several language models, with the corpus, the model, and the prompt all indepen-

dently configurable—and thus a faithful instance of the configuration problem CAFE targets.

A representative study defines factors over the base model (several open and hosted models), the retrieval corpus (a general legal corpus versus domain-specific collections), and the system prompt, evaluated on a curated catalogue of works-council questions each accompanied by an expert reference answer. Answers are compared against the reference on a three-level rubric, rated both by an LLM judge and by domain experts from partner labour organisations, who contribute through invitation-managed organisation accounts while the operating team administers studies. This deployment exercises every part of the system—factor registration, design generation, replicated execution, dual rating, and the analysis dashboard—on a fielded application with genuine domain-expert raters, the setting the planned evaluation (Section 6) will quantify.

6 Planned Evaluation

The following describes how we intend to validate CAFE; these are plans and design sketches, not yet results. The guiding principle is that “the platform reported configuration C as best” is a *demonstration*, whereas an *evaluation* must show that when the platform says “best,” it is right and the finding generalises. We plan four complementary studies, run primarily on public datasets (e.g. LegalBench for a legal flavour, plus open-domain QA such as Natural Questions/HotpotQA and a RAG-specific benchmark) so that validation does not depend on the still-growing catalogue of the deployed assistant, which is retained as the real-world demonstration.

- **Recovery (construct validity).** Build pipelines with planted ground truth—a strong vs. a deliberately weakened model, a correct vs. a sabotaged retrieval corpus, a sound vs. a degraded prompt—and verify that CAFE attributes variance to the factors we know matter, with the correct sign, and ranks the known-good configuration first. Because we control the truth, this needs no perfect dataset and directly tests whether the instrument is trustworthy.
- **Holdout / cross-validation (external validity).** Split *questions* (not runs) into estimation and held-out sets; check that the configuration ranking CAFE infers on the estimation set

predicts performance on held-out questions (Spearman correlation of predicted vs. realised ranking), using k -fold cross-validation when the question set is small. Where gold answers exist, score the held-out set against gold to fold in criterion validity.

- **Judge–human reliability.** Report Krippendorff’s α between the LLM judge and expert humans on a shared subset, establishing whether the automatic ratings that feed the analysis track expert judgement (Krippendorff, 2004; Haldar et al., 2025).
- **Reproducing known effects, and statistical power.** Check that CAFE recovers independently reported phenomena—e.g. large prompt-format sensitivity (Sclar et al., 2024) and non-monotonic effects of additional model calls (Chen et al., 2024)—and quantify the efficiency gain of the fractional-factorial plus ordinal-mixed-model pipeline over the prevailing single-run, interval-mean baseline (fewer runs to reach the same conclusion; effects detected that the baseline misses).

Together, recovery and holdout cover internal and external validity respectively, the reliability study validates the measurement, and the reproduction/power study situates the results against known findings; this is the evidence we will report in the camera-ready version.

7 Comparison to Existing Systems

Table 1 positions CAFE against representative evaluation systems. The distinguishing combination is not any single column—component evaluators score parts, optimizers search configurations, platforms run judges—but the integration of a designed experiment, ordinal variance attribution, dual human/judge rating with reliability, and explicit replication for nondeterminism, delivered as a deployed tool.

8 Conclusion

CAFE makes design of experiments the default for compound AI evaluation. By turning pipeline techniques into factors, generating factorial designs, replicating runs to handle nondeterminism, collecting ordinal judgments from both an LLM judge and human experts, and attributing variance with cumulative link mixed models, it is designed to let practitioners answer a question current tools leave open: which component drives answer qual-

System	Pipeline-level	DoE / factorial	Variance attrib.	Ordinal stats	Human + judge	Replication	Config. search
HELM (Liang et al., 2023)	–	–	–	–	–	–	–
RAGAS (Es et al., 2024)	part	–	–	–	–	–	–
ARES (Saad-Falcon et al., 2024)	part	–	–	–	–	–	–
DSPy (Khatab et al., 2024)	✓	–	–	–	–	–	✓
RAGSmith (Kartal et al., 2025)	✓	–	–	–	–	–	✓
LangSmith (LangChain, 2024)	✓	–	–	–	✓	–	–
Inspect AI (UK AI Security Institute, 2024)	✓	–	–	–	part	part	–
CAFE	✓	✓	✓	✓	✓	✓	–

Table 1: CAFE versus representative systems. “part” denotes partial support (e.g. component-only scoring, or annotation without designed experiments). CAFE’s distinguishing trait is not any single column but the combination of designed experiment, ordinal variance attribution, dual human/judge rating, and replication for nondeterminism; conversely, optimizers such as DSPy and RAGSmith provide the large-scale *configuration search* that CAFE deliberately leaves to future work.

ity, by how much, and how reliably we know it. The platform is deployed on a real legal-assistant application with domain-expert raters, and is open source. Our immediate next step is the validation programme of Section 6; longer term, we aim to add Bayesian-optimization-based refinement once screening has identified the factors worth tuning.

Availability

CAFE is released as open-source software under the Apache-2.0 license, with a live demonstration and a screencast; URLs are provided in the camera-ready version.

Limitations

CAFE is presented as a system with a validation plan rather than completed quantitative results; the studies of Section 6 are not yet run, and the present paper should be read accordingly. The statistical pipeline inherits the practical difficulties of mixed-effects modelling—convergence on small or sparse designs, and sensitivity to the random-effects specification—which the platform should surface to users rather than hide. Fractional factorial designs assume that high-order interactions are negligible; when that assumption fails, aliased effects can mislead, so the alias structure must be reported and read. The black-box treatment of components is deliberate but limiting: CAFE attributes variance to a technique as a whole, not to mechanisms inside it. The deployment covers a single domain (German labour law) and language; generality across domains is a claim the planned public-dataset studies are meant to support, not yet establish. Finally, an LLM-judge signal can be gamed: a configuration may be ranked highly because it suits the judge rather than users, which is precisely why judge–human reliability is reported

as a first-class output and why we advocate keeping humans in the loop.

Ethics Statement

The platform collects quality judgments from human domain experts; in the deployment these are members of partner labour organisations, who participate with informed consent and appropriate arrangement, and whose ratings are used only to evaluate system quality. Evaluation questions concern legal and works-council topics and contain no personal data. The application under evaluation provides legal *research support* and not legal advice; an evaluation platform such as CAFE is intended to make the quality and limitations of such systems measurable and transparent, which we regard as supporting their responsible deployment. We note the dual-use risk that an evaluation oriented toward an automatic judge could optimise systems toward that judge rather than toward people, and we mitigate this by foregrounding human-expert ratings and reporting judge–human agreement.

References

- Argilla. 2024. [Argilla: The collaboration platform for AI engineers and domain experts](#). Software.
- Berk Atil and 1 others. 2024. [Non-determinism of “deterministic” LLM settings](#). *arXiv preprint arXiv:2408.04667*.
- Matthew Barker and Andrew Bell. 2025. [Faster, cheaper, better: Multi-objective hyperparameter optimization for LLM and RAG systems](#). *arXiv preprint arXiv:2502.18635*.
- Pranav Bhagat, Sandeep Shastri, and 1 others. 2025. [Evaluating compound AI systems through behaviors, not benchmarks](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*.

- Brynjolfur Bjarnason and 1 others. 2026. [On randomness in agentic evals](#). *arXiv preprint arXiv:2602.07150*.
- Pierre Boyeau, Anastasios N. Angelopoulos, Nir Yosef, Jitendra Malik, and Michael I. Jordan. 2025. [AutoEval done right: Using synthetic data for model evaluation](#). In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*.
- Jiayi Chen and 1 others. 2025a. [From standalone LLMs to integrated intelligence: A survey of compound AI systems](#). *arXiv preprint arXiv:2506.04565*.
- Lingjiao Chen, Jared Quincy Davis, Boris Hanin, Peter Bailis, Ion Stoica, Matei Zaharia, and James Zou. 2024. [Are more LLM calls all you need? towards scaling laws of compound inference systems](#). *arXiv preprint arXiv:2403.02419*.
- Lingjiao Chen and 1 others. 2025b. [Optimizing model selection for compound AI systems](#). *arXiv preprint arXiv:2502.14815*.
- Yann Dubois and 1 others. 2025. [Skewed score: A statistical framework to assess autograders](#). *arXiv preprint arXiv:2507.03772*.
- Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. 2024. [RAGAs: Automated evaluation of retrieval augmented generation](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (EACL): System Demonstrations*.
- Adam Fisch, Joshua Maynez, R. Alex Hofer, Bhuwan Dhingra, Amir Globerson, and William W. Cohen. 2024. [Stratified prediction-powered inference for effective hybrid evaluation of language models](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Jiawei Gu and 1 others. 2024. [A survey on LLM-as-a-judge](#). *arXiv preprint arXiv:2411.15594*.
- Luke Guerdan and 1 others. 2025. [Validating LLM-as-a-judge systems under rating indeterminacy](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Jennifer Haase and 1 others. 2026. [Within-model vs between-prompt variability in large language models for creative tasks](#). *arXiv preprint arXiv:2601.21339*.
- Rajarshi Halder and 1 others. 2025. [Rating roulette: Self-inconsistency in LLM-as-a-judge frameworks](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*.
- Steve Han and 1 others. 2025. [Judge’s verdict: A comprehensive analysis of LLM judge capability through human agreement](#). *arXiv preprint arXiv:2510.09738*.
- David M. Howcroft and Verena Rieser. 2021. [What happens if you treat ordinal ratings as interval data? human evaluations in NLP are even more underpowered than you think](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Muhammed Yusuf Kartal and 1 others. 2025. [RAG-Smith: A framework for finding the optimal composition of retrieval-augmented generation methods across datasets](#). *arXiv preprint arXiv:2511.01386*.
- Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. 2024. [DSPy: Compiling declarative language model calls into self-improving pipelines](#). In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Klaus Krippendorff. 2004. *Content Analysis: An Introduction to Its Methodology*, 2nd edition. Sage.
- LangChain. 2024. [LangSmith: A platform for LLM application evaluation and observability](#). Software.
- Yu-Ang Lee and 1 others. 2025. [Compound AI systems optimization: A survey of methods, challenges, and future directions](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Percy Liang, Rishi Bommasani, Tony Lee, and 1 others. 2023. [Holistic evaluation of language models](#). *Transactions on Machine Learning Research (TMLR)*.
- Lennart Luetzgau and 1 others. 2025. [HiBayES: A hierarchical bayesian modeling framework for AI evaluation statistics](#). *arXiv preprint arXiv:2505.05602*.
- Fabian Lukassen and 1 others. 2026. [From XAI to stories: A factorial study of LLM-generated explanation quality](#). *arXiv preprint arXiv:2601.02224*.
- Solomon Messing. 2026. [Hidden measurement error in LLM pipelines distorts annotation, evaluation, and benchmarking](#). *arXiv preprint arXiv:2604.11581*.
- Zairah Mustahsan and 1 others. 2025. [Stochasticity in agentic evaluations: Quantifying inconsistency with intraclass correlation](#). *arXiv preprint arXiv:2512.06710*.
- Sangwoo Park and 1 others. 2025. [Adaptive prediction-powered AutoEval with reliability and efficiency guarantees](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Jon Saad-Falcon, Omar Khattab, Christopher Potts, and Matei Zaharia. 2024. [ARES: An automated evaluation framework for retrieval-augmented generation systems](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.

- Andrea Saltelli, Marco Ratto, Terry Andres, Francesca Campolongo, Jessica Cariboni, Debora Gatelli, Michaela Saisana, and Stefano Tarantola. 2008. *Global Sensitivity Analysis: The Primer*. John Wiley & Sons.
- Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. Quantifying language models’ sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Bani Syiem and Eduardo Velloso. 2026. Better assumptions, stronger conclusions: The case for ordinal regression in HCI. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems (CHI)*.
- Jack E. Taylor, Guillaume A. Rousselet, Christoph Scheepers, and Sara C. Sereno. 2023. Rating norms should be calculated from cumulative link mixed effects models. *Behavior Research Methods*.
- UK AI Security Institute. 2024. *Inspect AI: An open-source framework for large language model evaluations*. Software.
- Alan R. Vazquez and 1 others. 2025. A systematic assessment of large language models for constructing two-level fractional factorial designs. *arXiv preprint arXiv:2512.17113*.
- Bo Wang and 1 others. 2025. On the reproducibility limitations of RAG systems. *arXiv preprint arXiv:2509.18869*.
- Shirley Wu, Parth Sarthi, and 1 others. 2026. Optimas: Optimizing compound AI systems with globally aligned local rewards. In *The Fourteenth International Conference on Learning Representations (ICLR)*.
- Mert Yuksekgonul, Federico Bianchi, Joseph Boen, Sheng Liu, Pan Lu, Zhi Huang, Carlos Guestrin, and James Zou. 2025. Optimizing generative AI by backpropagating language model feedback. *Nature*, 639. TextGrad.
- Matei Zaharia, Omar Khattab, Lingjiao Chen, Jared Quincy Davis, Heather Miller, Chris Potts, and James Zou. 2024. The shift from models to compound AI systems. Berkeley Artificial Intelligence Research (BAIR) Blog.

A User-Interface Walkthrough

This appendix illustrates the experimenter-facing interface referenced in Section 3. Figures 2–5 follow a single study from technique inspection through design definition, rating, and analysis.




Figure 2: The *Techniques* page lists every registered stage with its techniques (the levels of each factor), so the testable configuration space is exactly what the codebase implements.




Figure 3: Defining a study: the experimenter selects questions, ticks the levels of each factor to vary, and chooses a design; CAFE reports the resulting number of cells and runs.

DIVA Eval

Mitwirken

Fragen

Bewerten

test@test.de

Test Org, delte later

Sign out

Bewertung

Noch ca. 978 Antworten zu bewerten

FRAGE

Darf der Betriebsrat in Personalentscheidungen direkt eingreifen?

REFERENZANTWORT

Kein Vetorecht gegen Einstellungen/Versetzungen an sich Aber: Zustimmungsverweigerung bei klaren gesetzlichen Gründen möglich Keine Auswahlentscheidung, nur Prüfung der Rechtmäßigkeit. Bei Kündigungen: Der Betriebsrat ist vor jeder Kündigung zu hören. Der Arbeitgeber hat ihm die Gründe für die Kündigung mitzuteilen. Eine ohne Anhörung des Betriebsrats ausgesprochene Kündigung ist unwirksam. (BetrVG § 102)Die Konsultation mit dem Betriebsrat ist bei einer Massenentlassung nur dann ordnungsgemäß, wenn in einem faktisch abhängigen Unternehmen nicht nur die vordergründigen Gründe, sondern die Hintergründe ausführlich mitgeteilt werden. LAG Berlin-Brandenburg, Urteil vom 26.11.2015

ANTWORT

[Platzhalter-Antwort, Konfiguration: GPT-OSS-120B + Verdi RAG]

Zur Frage »Darf der Betriebsrat in Personalentscheidungen direkt eingreifen?...« wäre eine ausführliche, paragraphenfeste Antwort zu geben. Dieser Text dient als Platzhalter, damit das Bewertungs-Frontend vor Anbindung der echten DIVA-Pipeline getestet werden kann. Ersetze diese Antworten durch echte DIVA-Outputs, sobald die Run-Komponente bereit ist.

WIE SCHNEIDET DIESE ANTWORT GEGENÜBER DER REFERENZANTWORT AB?

schlechter

WORSE

ebenbürtig

ON_PAR

besser

BETTER

BEGRÜNDUNG (OPTIONAL)

Welche Aspekte führen zu deiner Einschätzung?

Überspringen

Bewertung absenden

Figure 4: The human-rating interface presents an answer alongside the ordinal rubric; the same rubric is applied by the configurable LLM judge, enabling judge–human reliability analysis.




Figure 5: The results dashboard reports each factor’s relative contribution, the best configuration, the verdict distribution, and pipeline-stage timings.